

Queries and Time

Flink SQL Training

<https://github.com/ververica/sql-training>

Temporal Conditions are Common in Stream Processing

- Smoothing and aggregating data based on time
 - *"Compute average over the last 1 minute"*
- Enriching streaming data with data from other sources
 - *"Join with most recent exchange rate"*
- Stream monitoring, pattern matching, and alerting
 - *"Emit an alert if 3 unsuccessful attempts occurred within 5 minutes"*
- Common types of data
 - User interactions: clicks, (mobile) apps
 - Logs: applications, machines, network hardware
 - Transactions: credit cards, ad serving
 - Sensors: mobile phones, cars, IoT



Characteristics of Queries with Temporal Conditions

- Input tables are append-only. Rows are not updated after they were inserted.
 - Table schema often defines an event-time (or processing-time) attribute
- The query consists of row-at-a-time and temporal operators
 - Filter, Projection, Windowed aggregations, Interval join, Temporal-table join, Pattern matching
 - “Traditional” stream processing operations
 - Temporal operators track time progress and compute final results
- Query output is append-only. Emitted result rows are never updated



Event & Processing Time Attributes

How Is Time Handled in Flink SQL?

- Flink SQL supports Event and Processing Time
- Tables may include *Time Attributes*
 - Time Attributes provide access to event or processing time
 - Time Attributes are (mostly) treated as regular attributes
 - Time Attributes have special type extended from SQL TIMESTAMP
- Time Attributes are declared with the table schema



Event Time Attribute

- Event time attribute is of type `TIMESTAMP(3)` and carries an actual timestamp
- Watermarks are meta-records and generated based on observed timestamps
 - Operators use watermarks to track the progress of time and reason about data completeness
- Event time attribute can be used like a regular `TIMESTAMP` attribute
 - Loses its event-time property when it is modified or when operators change the order of rows

```
CREATE TABLE clicks (  
  user VARCHAR,  
  url VARCHAR,  
  cTime TIMESTAMP(3),  
  WATERMARK FOR cTime AS cTime - INTERVAL '2' MINUTE  
)
```



Processing Time Attribute

- Processing time attributes are virtual and do not hold data
 - Local machine time is queried when the attribute is accessed
- A processing time attribute can be used like a regular **TIMESTAMP**
 - Loses its processing time property when it is modified

```
CREATE TABLE clicks (  
  user VARCHAR,  
  url VARCHAR,  
  cTime AS PROCTIME() // is of type TIMESTAMP(3)  
)
```



Temporal Operators

Temporal Operators

- Temporal operators associate records with each other based on temporal conditions
 - GROUP BY window aggregation
 - OVER window aggregation
 - Time-windowed join
 - Join with a temporal table (enrichment join)
 - Pattern matching (MATCH_RECOGNIZE)
- } Discussed in a later session
- Operators track progress in time to decide when input is complete
 - Operators emit final result rows that will not be updated
 - Operators discard state (records and results) that are no longer needed
 - Temporal operators require time attributes in specific clauses



Temporal Aggregation

Temporal Aggregation

- Flink SQL supports two types of temporal aggregations
 - GROUP BY window aggregation
 - OVER window aggregation
- The clicks table is used for the following examples
 - cTime (clickTime) is an event time attribute.

user	cTime	url
Mary	12:00:00	./home
Bob	12:00:00	./cart
Mary	12:00:05	./prod?id=1
...



GROUP BY Window Aggregation

Compute number of clicks per hour and user

```
SELECT user,  
       TUMBLE_END(cTime, INTERVAL '1' HOURS) AS endT,  
       COUNT(url) AS cnt  
FROM clicks  
GROUP BY TUMBLE(cTime, INTERVAL '1' HOURS),  
       user
```

Time attribute

Time attribute



GROUP BY Window Aggregation

clicks		
user	cTime	url
Mary	12:00:00	./home
Bob	12:00:00	./cart
Mary	12:02:00	./prod?id=2
Mary	12:55:00	./home
Bob	13:01:00	./prod?id=4
Liz	13:30:00	./cart
Liz	13:59:00	./home
Mary	14:00:00	./prod?id=1
Liz	14:02:00	./prod?id=8
Bob	14:30:00	./prod?id=7
Bob	14:40:00	./home

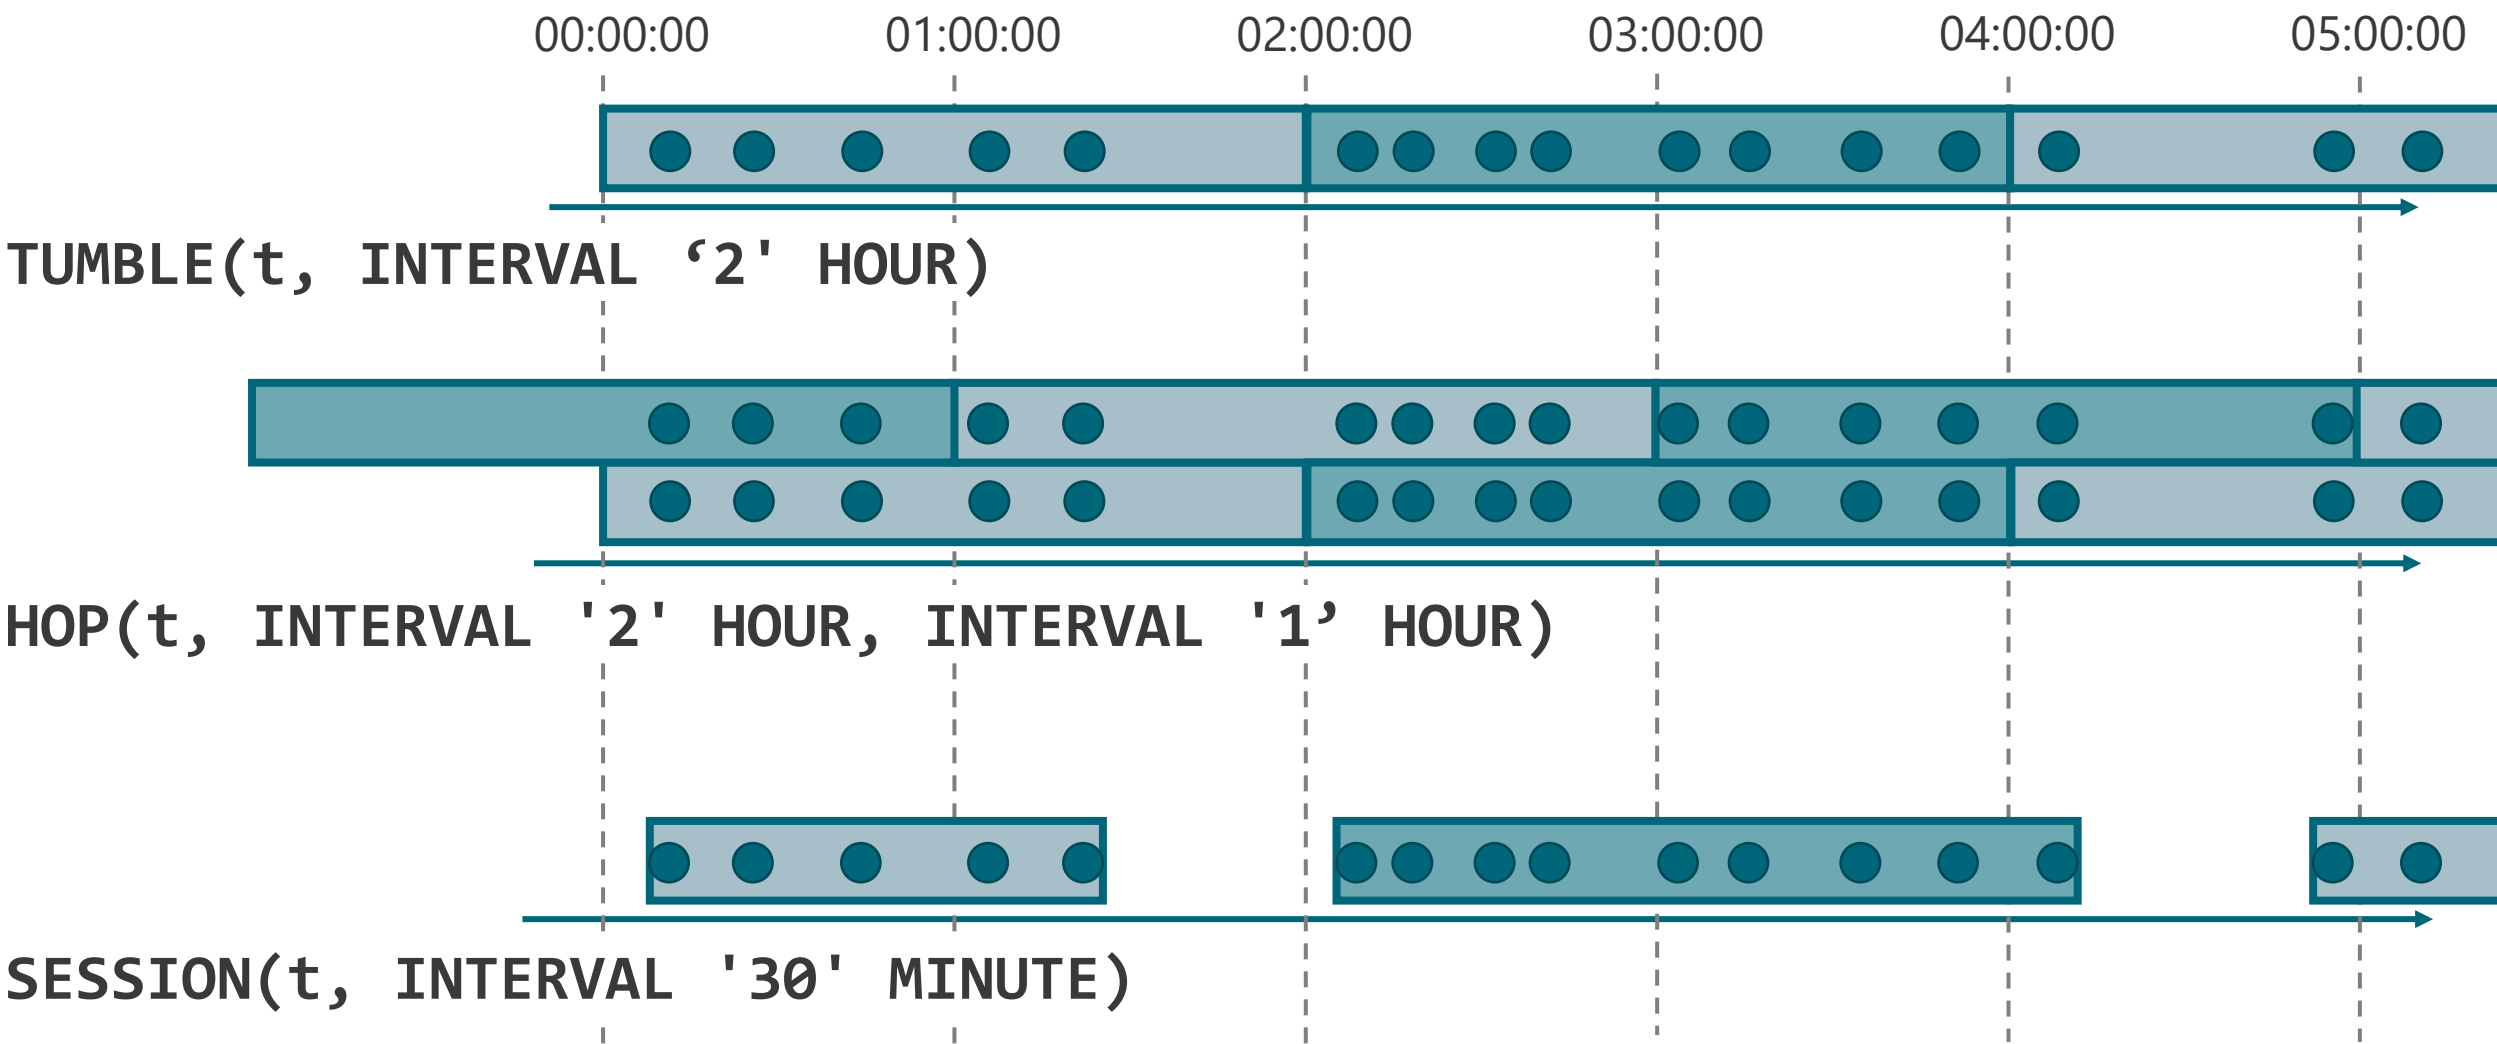
```
SELECT
  user,
  TUMBLE_END(
    cTime,
    INTERVAL '1' HOURS)
  AS endT,
  COUNT(url) AS cnt
FROM clicks
GROUP BY
  user,
  TUMBLE(
    cTime,
    INTERVAL '1' HOURS)
```

result		
user	endT	cnt
Mary	13:00:00	3
Bob	13:00:00	1
Bob	14:00:00	1
Liz	14:00:00	2
Mary	15:00:00	1
Bob	15:00:00	2
Liz	15:00:00	1

Rows are appended to result table.




GROUP BY Windows



OVER Window Aggregation

Compute for each click
how often its URL was clicked in the previous 2 hours

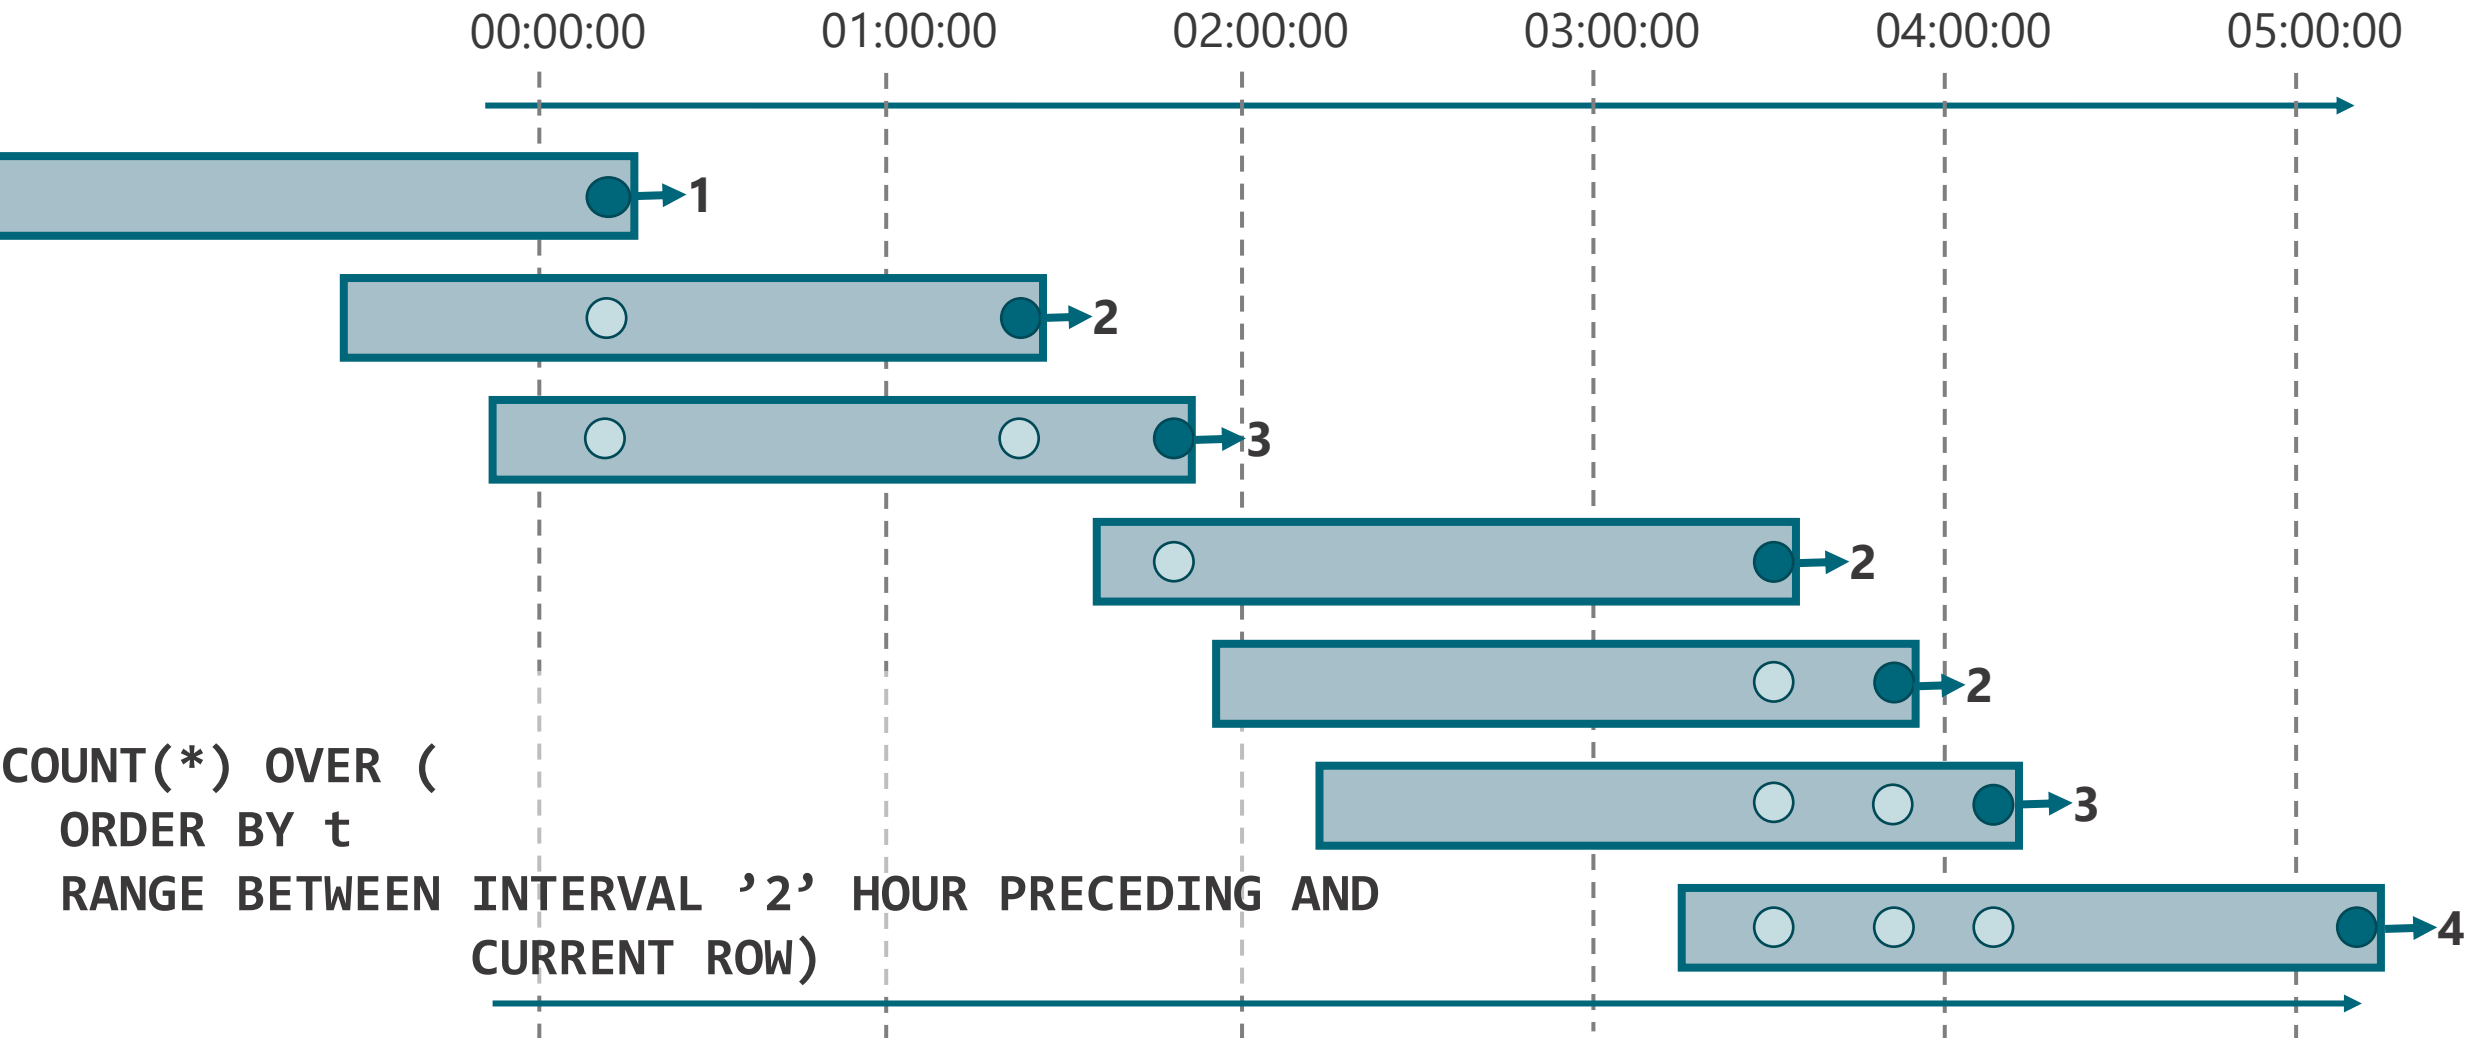
```
SELECT
  user,
  url,
  COUNT(*) OVER w
FROM clicks
WINDOW w AS (
  PARTITION BY url
  ORDER BY cTime
  RANGE BETWEEN INTERVAL '2' HOUR PRECEDING AND CURRENT ROW)
```



Time attribute



OVER Windows



Validity of Time Attributes

Validity of Event-Time Attributes

- Attributes lose their event-time property when watermark alignment cannot be guaranteed anymore
 - Event-time attributes become regular `TIMESTAMPS`
 - Cannot be used in temporal operations anymore

- Modifying a time attribute

```
SELECT FLOOR(cTime TO MINUTE) FROM clicks
```

- Operators that do not preserve the timestamp order in their output
 - Joins without a temporal join condition (discussed later in detail)
 - Non-windowed aggregations

```
SELECT cTime, COUNT(*) FROM clicks GROUP BY cTime
```



Validity of Processing-Time Attributes

- Attributes lose their processing-time property when they are modified
 - A processing-time attribute is materialized when it is accessed and becomes regular a TIMESTAMP
 - Cannot be used in temporal operations anymore

```
SELECT FLOOR(cTime TO MINUTE) FROM clicks
```



Summary

Summary

- Many streaming SQL use cases require temporal operators
 - Input and output are append-only tables
 - Queries use record-at-a-time and temporal operators
- Time attributes are defined in the table schema
 - Event-time and/or processing-time attributes are supported
- Temporal operations associate input rows based on their time attributes
 - Produce final results, i.e., emitted rows are never updated
 - Automatically manage their state



Hands On Exercises

Queries and Time

Continue with the hands-on exercises in
"Queries and Time"

<https://github.com/ververica/sql-training/wiki/Queries-and-Time>

We are here to help!





ververica

www.ververica.com

@VervericaData